

SHIVA GUPTA

+91 7877993048

shivafy02@gmail.com

LinkedIn

GitHub

LeetCode

EDUCATION

B.Tech, Computer Science and Engineering, Manipal University Jaipur, 7.7 CGPA

2023 - 2027

EXPERIENCE (6 MONTHS)

Reliquis Consulting

New Delhi, India

Full Stack Intern

June 2025 – August 2025

- Designed and implemented a production-ready software architecture for designmumbai.com, as a junior developer optimizing responsive UI layouts by using Redux and implementing CI/CD pipelines to streamline development and deployment workflows.
- Optimized application performance, minimized database load, and enhanced overall reliability by architectural integration of Redis-based distributed lock system so frequent database calls reduce.

Frontend Design Intern

June 2024 – August 2024

- Developed a responsive, high-performance e-commerce platform from Figma UI/UX prototypes, ensuring high-fidelity visual execution, structural SEO, and smooth animations using Bootstrap, JavaScript and libraries.
- Leveraged strict Git/GitHub version control workflows for agile team collaboration, resolved merge conflicts, and implemented cross-browser performance optimization best practices to reduce page load times.

PROJECTS

Brainly [ReactJS, MongoDB, TypeScript, Langchain, Docker, AWS(EC2), CI/CD, Vercel]

[GitHub](#)

- Built a Smart Search Engine for Notes: Created a platform that lets users search and retrieve their digital notes based on the meaning of their text rather than just exact keywords.[**Langchain and MongoDB Vector Indexing**]
- Automated the Launch and Update Process: Set up a modern system that automatically builds, tests, and deploys the app whenever changes are made so i don't have to take it down and relaunch to cause downtime of website.[**AWS (EC2) and GitHub actions**]
- Created Secure Multi-User Sharing: Designed the platform to handle multiple users safely, 'share brain' feature allowing people to share workspaces in read mode without compromising their data from being altered.[**JWT and permission-controlled access**]

Legal intelligence Agent [LangGraph, Langsmith, BeautifulSoup4, bautasaurus, Multi-stage docker, Cloud run]

[GitHub](#)

- Built a Self-Diagnostic and corrective Multi agent orchestration: Created smart legal document analysis agentic orchestration that automatically detects its own errors during research, diagnoses why it failed, and tries a different approach to fix itself on the fly, and gracefully aborts when no solution is possible instead of crashing the system.[**LangGraph and BeautifulSoup4**]
- Made the System Ultra-Reliable and Reusable: Stripped out complex error-handling code from the core AI logic by creating a single, centralized manager that safely runs all 6 legal tools and handles bugs behind the scenes, so essentially a single node spawns all agents rather than making a redundant architecture with code repetition.[**Centralized ToolExecutor architecture**]
- Created a Live Tracking Dashboard: Built a visual workspace that shows exactly how the AI thinks in real-time, tracing sequence of node execution, every single decision the architecture makes, letting users see every tool attempt, failure diagnosis, and self-correction decision as it happens.[**Streamlit UI**]
- Monitoring latency at every node response using langsmith SDK reducing the response time by 3s, using a scraping library instead of browser search api key.[**Langsmith**]
- Launching the agent through container deployment and Cloud run(Google) this setup only uses the credits when being run.

CompIntel [LangGraph, BeautifulSoup4, httpx, Docker]

[GitHub](#)

- Built an Automated multi-agent Research Team: Created a system of specialized AI agents that automatically pull live market data, financials, and competitor info for any company by orchestrating workflows.[**LangGraph, Tavily and Alpha Vantage**]
- Made the AI Reliable and Bug-Free: Prevented the AI system from getting stuck in endless loops or crashing from bad data by locking down its decision-making.[**Pydantic enums and OpenAI structured outputs**]
- Created a Live Execution Dashboard: Built a highly visual, real-time user interface that lets users watch the AI agents work, think, and route data live, completely containerized using Docker for easy setup.[**Streamlit and Docker**]

Prompt compressor for LLMs [PyTorch, HuggingFace Transformers, DistilBERT, Streamlit]

[GitHub](#)

- Built a learned prompt compression pipeline to reduce LLM token usage by fine tuning a DistilBERT-based token scoring model on 500 SQuAD 2.0 passages, the tokens were scored based on a 'surprisal' method.
- Reduced token count by up to 70 percent outperforming GPT-2 perplexity baselines under aggressive compression, significantly accelerating LLM inference speed and lowering API operational costs.

ACHIEVEMENTS

Startup Weekend Jaipur — Technical Lead

[Live Website](#)

- Served on the Executive Committee as Technical Lead and designed and developed the official event website, delivering a production-ready platform for registrations and engagement ([GitHub Repository](#)).